

# Foundations of DL

Deep Learning



ALF




Alfredo Canziani, Ritchie Ng  
@alfcnz, @RitchieNg


# Generative models

(Semi-) Unsupervised learning

# Super resolution








Garcia (2016) srez --- [github.com/david-gpu/srez](https://github.com/david-gpu/srez)



# Inpainting


VAE



GAN

Yeh (2017)  
Semantic Image Inpainting  
with Deep Generative Models

[bit.ly/DCGAN-inpainting](http://bit.ly/DCGAN-inpainting)





# Caption to image

This vibrant red bird has  
a pointed black beak

This bird is yellowish  
orange with black wings

The bright blue bird has  
a white coloured belly

Reed (2016) Generative  
adversarial text to image synthesis



[github.com/reedscot/icml2016](https://github.com/reedscot/icml2016)



# Autoencoders

Unsupervised learning / Generative models

# Autoencoder



$$h = f(W_h x + b_h)$$

$$\hat{x} = g(W_x h + b_x)$$

$$x, \hat{x} \in \mathbb{R}^n$$

$$h \in \mathbb{R}^d$$

$$W_h \in \mathbb{R}^{d \times n}$$

$$W_x \in \mathbb{R}^{n \times d}$$

If “tight weights”, then

$$W_x \doteq W_h^\top$$



# Reconstruction losses

$$\mathcal{L} = \frac{1}{m} \sum_{j=1}^m \ell(\mathbf{x}^{(j)}, \hat{\mathbf{x}}^{(j)})$$


binary input

$$\ell(\mathbf{x}, \hat{\mathbf{x}}) = - \sum_{i=1}^n [x_i \log(\hat{x}_i) + (1 - x_i) \log(1 - \hat{x}_i)]$$

real valued input

$$\ell(\mathbf{x}, \hat{\mathbf{x}}) = \frac{1}{2} \|\mathbf{x} - \hat{\mathbf{x}}\|^2$$


# Under-/over-complete hidden layer



$\hat{x}$


$h$

$x$



# Denoising autoencoder

$$\tilde{x} \sim p(\tilde{x} \mid x)$$



**We assume** we are injecting the **same noisy distribution** we're going to observe in reality. In this way, we can learn how to robustly recover from it.




# Contractive autoencoder


penalises insensitivity to  
reconstruction directions

penalises sensitivity  
to the **any** direction

$$\ell(\mathbf{x}, \hat{\mathbf{x}}) = \boxed{\ell_{\text{reconstruction}}} + \boxed{\lambda \|\nabla_{\mathbf{x}} \mathbf{h}\|^2}$$





## Basic auto-encoder



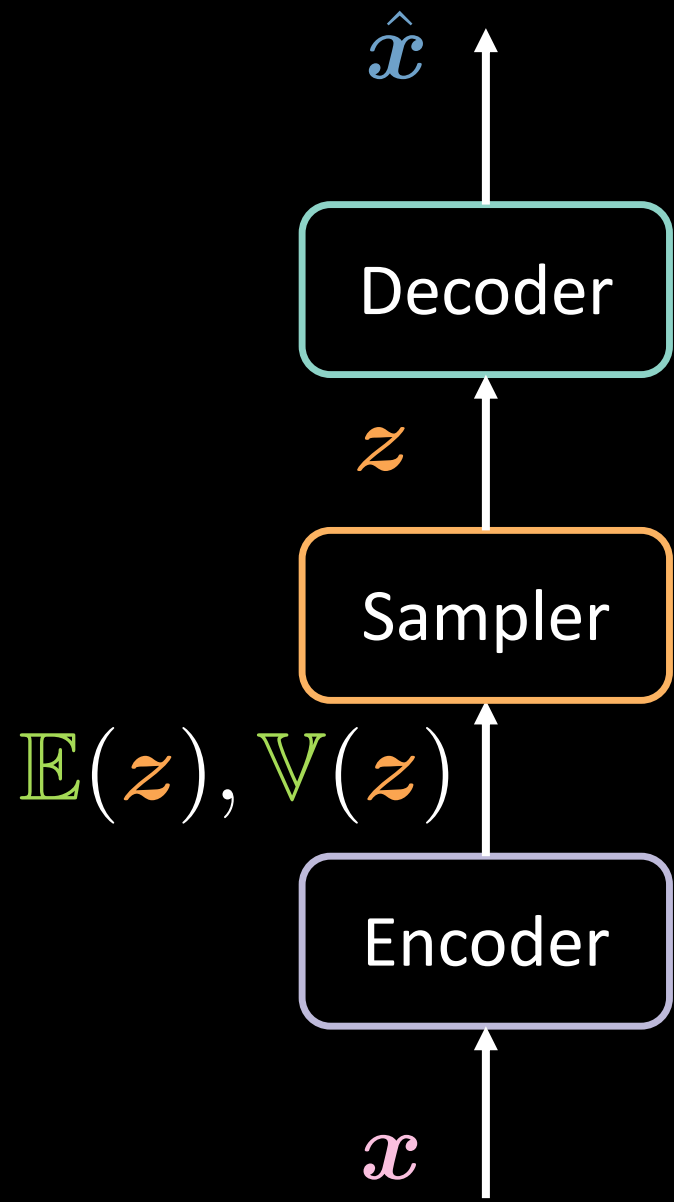
- How can we “go back”:
- learn the distribution
  - or
  - enforce some structure

# Auto-encoder (recap)






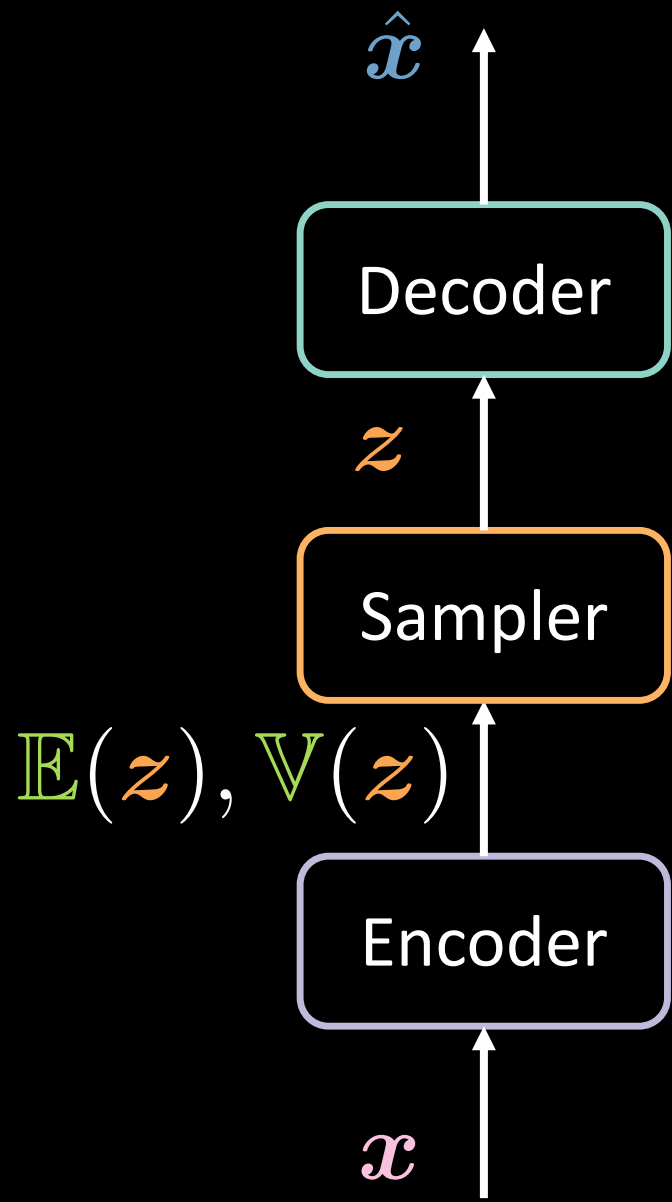
Variational auto-encoder



Classic auto-encoder




# Variational auto-encoder



$$\text{decoder} : \mathcal{Z} \rightarrow \mathbb{R}^n$$
$$z \mapsto \hat{x}$$

$$\text{encoder} : \mathcal{X} \rightarrow \mathbb{R}^{2d}$$
$$x \mapsto h$$

# Variational auto-encoder



To go from the latent to the input space we need to:

- learn the distribution
- or
- enforce some structure




# Variational auto-encoder


$$\ell(\mathbf{x}, \hat{\mathbf{x}}) = \ell_{\text{reconstruction}} + \beta \ell_{\text{KL}}(\mathbf{z}, \mathcal{N}(\mathbf{0}, \mathbb{I}_d))$$

$$= \ell_{\text{reconstruction}} +$$

$$+ \beta \sum_{i=1}^d \left( \underbrace{\mathbb{V}(\mathbf{z}) - \log [\mathbb{V}(\mathbf{z})] - \mathbf{1}}_{\mathbf{v}} + \mathbb{E}(\mathbf{z})^2 \right)_i$$



$\mathbb{V}(\mathbf{z})_i$





# Variational auto-encoder


$$\ell(\mathbf{x}, \hat{\mathbf{x}}) = \ell_{\text{reconstruction}} + \beta \ell_{\text{KL}}(\mathbf{z}, \mathcal{N}(\mathbf{0}, \mathbb{I}_d))$$

$$= \ell_{\text{reconstruction}} +$$

$$+ \beta \sum_{i=1}^d \left( \underbrace{\mathbb{V}(\mathbf{z}) - \log [\mathbb{V}(\mathbf{z})] - \mathbf{1}}_{\mathbf{v}} + \overbrace{\mathbb{E}(\mathbf{z})^2}^{\mathbf{u}} \right)_i$$



$\mathbb{V}(\mathbf{z})_i$

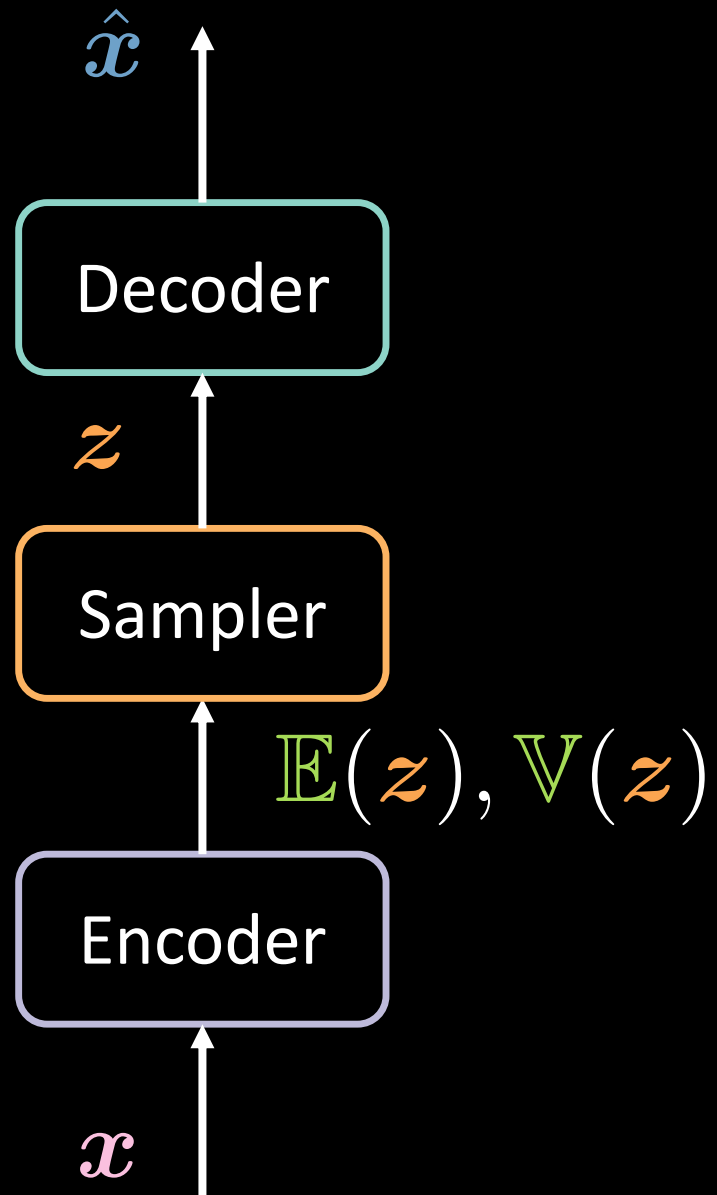


# Generative adversarial nets


Unsupervised learning / Generative models



Variational auto-encoder



Generative adversarial network



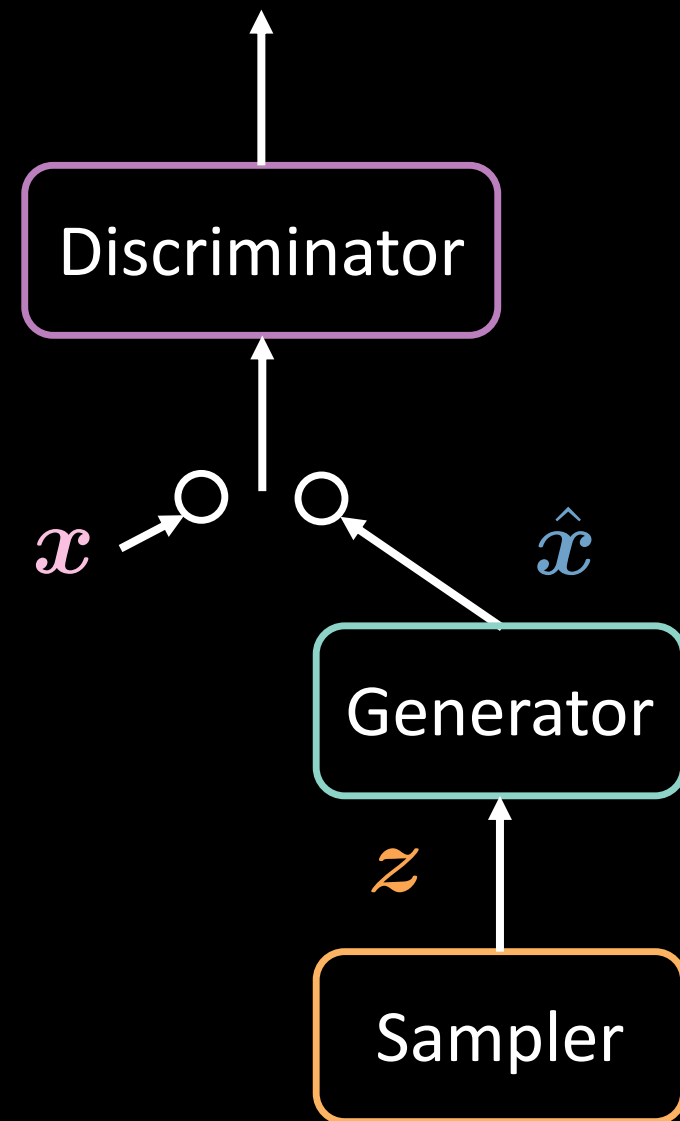
$$D : \mathbb{R}^n \rightarrow (0, 1)$$

$$x \vee \hat{x} \mapsto \ell$$


$$G : \mathcal{Z} \rightarrow \mathbb{R}^n$$

$$z \mapsto \hat{x}$$

Generative adversarial network



# Generative adversarial network



# Value function

$$V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \\ + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D[G(\mathbf{z})])]$$

$$\min_G \max_D V(D, G)$$